

A crossed random-effects latent class model for annotation ensembles of large language models

Author information withheld for double-blind review^{1*}

^{1*} Affiliation withheld.

Abstract

Large language models are increasingly used to label text at scale, but every label is produced by a *configuration*—a model, a prompt and a stochastic decoding run—and configurations that share training corpora and alignment procedures do not err independently. Classical aggregation models, from Dawid and Skene onward, assume conditional independence of raters given the true class. We propose a latent class model whose dependence structure follows the factorial design of the annotation ensemble: item-level random effects shared by all configurations, crossed with item-by-model and item-by-prompt effects, with decoding runs exchangeable. Estimated by Bayesian inference with the latent class marginalised and the shared item effect integrated by quadrature, the model yields point estimates without a gold standard, with identification supported empirically by simulation, and decomposes annotation error variance into item, model, prompt and run components. In a simulation study, Dawid–Skene increasingly overstates annotator accuracy and biases prevalence as dependence grows; the proposed model reduces the prevalence bias several-fold (from **+0.05** and **+0.10** to about **+0.01** and **+0.04** at moderate and strong dependence), restores nominal interval coverage and yields better-calibrated posterior class probabilities. In an application to policy-topic coding of 800 US congressional bills by three model families, five prompts and three runs, the ensemble is near-unanimous and all aggregators reach similar accuracy and calibration against human codes; the gain lies in the prevalence, whose bias is halved (from **+0.031** to **+0.015**), and in the per-configuration accuracy estimates. The variance decomposition guides how many models, prompts and runs to use.

Keywords: latent class analysis, conditional dependence, random effects, large language models, text annotation, measurement error

MSC Classification: 62H30 , 62F15 , 62J12 , 62P25 , 68T50

1 Introduction

Text classification by large language models has moved, within three years, from a curiosity to routine research infrastructure. Studies in the social sciences, management and public administration now use LLMs to label sentiment, policy topics, stance, toxicity or compliance in corpora far too large for manual coding (Gilardi et al. 2023; Carlson and Burbano 2026). The attraction is obvious: labels arrive at a fraction of the cost and time of human annotation, and early comparisons suggested accuracy on a par with crowd workers.

The reliability of those labels is a different matter. An LLM label is not the output of a single, stable instrument. It is the output of a configuration: a particular model, a particular prompt wording, and a particular stochastic decoding run. Each of these varies, and each variation moves the label. Semantically equivalent prompts on the same task can shift accuracy by more than twenty points (Liu 2026); nominally deterministic decoding settings do not in fact return deterministic output (Atil et al. 2025); and re-running the same pipeline over several months yields markedly different labels, in part because API-served models are silently updated or retired (Barrie et al. 2025). At scale, these choices have downstream consequences. Baumann et al. (2025) replicated 37 annotation tasks from 21 published studies across 18 models and found that the choice of model and prompt alone was enough to change the sign or significance of substantive conclusions—a phenomenon they call LLM hacking. Camuffo et al. (2026) report that in their strategy-research setting the choice of model determined the outcome in 90% of annotations of identical business-model comparisons, with cross-model agreement barely above chance (mean Cohen’s $\kappa = 0.026$).

The natural statistical response is to treat each configuration as a rater and aggregate. This is precisely what recent methodological guidance recommends: query several models with several prompts, replicate, and combine the labels with an aggregation model such as Dawid and Skene (1979) or GLAD (Whitehill et al. 2009) that estimates each rater’s error rates jointly with the latent true labels (Camuffo et al. 2026). The recommendation is sensible, but it imports an assumption that the LLM setting violates in a specific and predictable way. Dawid–Skene and its descendants assume that raters are conditionally independent given the true class. Human crowd workers approximate this reasonably well; they err for idiosyncratic reasons. LLM configurations do not. Two prompts sent to the same model share that model’s training data, tokenisation and alignment; two models given the same prompt share the prompt’s framing and any ambiguity in the item itself; and all configurations share whatever makes a particular text genuinely hard to classify. The errors are therefore correlated, and the correlation has structure.

The diagnostic-testing literature met exactly this problem thirty years ago, when multiple imperfect tests applied to the same patient were found to share unmeasured patient-level characteristics. The conditional-independence latent class model fits such data poorly (Qu et al. 1996) and overstates the tests’ sensitivities and specificities (Vacek 1985). Qu et al. (1996) introduced subject-level random effects into the latent class model to absorb the shared dependence, and the approach has since been extended, tested and implemented in software (Beath 2017; Xu et al. 2019); methods for detecting the local dependence it addresses have also been evaluated (Oberski

et al. 2013). Within the classification community, relaxations of local independence have been developed for model-based clustering of categorical data (Marbac et al. 2016; Oberski 2016). What these models do not do is exploit the fact that, in an LLM ensemble, the rater is itself a designed object with a known factorial structure.

This paper makes three contributions. First, we propose a latent class model—the crossed random-effects latent class model, CRE-LCM—in which the dependence between LLM configurations is neither ignored nor left unrestricted but derived from the design of the annotation ensemble: an item effect shared by all configurations, crossed with item-by-model and item-by-prompt effects, with repeated decoding runs treated as exchangeable given these effects. Dawid–Skene and the one-factor model of Qu et al. (1996) are nested special cases. Second, we study identifiability. The crossed effects are identified by the replication structure of the design; the global item effect competes with the latent class itself, and we characterise the conditions and the remedies under which the model yields point estimates, with identification supported empirically by the simulation study rather than proved, in contrast to the partial-identification route recently proposed for the same problem (Chen et al. 2026). Third, we show through simulation that independence-based aggregation overstates accuracy and biases prevalence monotonically in the strength of dependence, that CRE-LCM reduces the prevalence bias several-fold (from +0.05 to +0.01 at moderate and from +0.10 to +0.04 at strong dependence) and restores nominal coverage of its credible intervals, and that its variance decomposition provides direct guidance for annotation design: whether to add models, prompts or runs. An application to policy-topic coding of legislative texts, with human reference labels held out for validation, illustrates the method.

The remainder of the paper is organised as follows. Section 2 reviews related work. Section 3 presents the model, its special cases, identifiability and estimation. Section 4 reports the simulation study. Section 5 presents the application. Section 6 discusses design implications and limitations.

2 Background and related work

2.1 Aggregating labels from multiple raters without a gold standard

The joint estimation of true labels and rater error rates was formalised by Dawid and Skene (1979), who modelled each rater by a class-conditional confusion matrix and used EM. The model remains the reference point for crowdsourced labelling and unsupervised ensemble aggregation. Extensions add item difficulty (GLAD; Whitehill et al. 2009), Bayesian priors and hierarchical structure for sparse-label regimes, spectral initialisation, and online estimation. Within the classification literature, Salter-Townshend and Murphy (2014) fitted a finite mixture to annotator bias matrices in order to cluster inexpert sentiment annotators, treating annotator heterogeneity itself as a clustering problem; Raykar and Yu (2012) ranked and filtered crowd annotators by estimated reliability. Agreement statistics such as multi-rater kappas (Warrens 2010) and weighted kappas (Warrens 2013) summarise concordance descriptively but

do not estimate accuracy against a latent truth. The common thread in the model-based approaches is conditional independence of raters given the true class. The current LLM-annotation guidance—treat each model or model-prompt pair as an annotator and apply Dawid-Skene or GLAD via standard packages (Camuffo et al. 2026)—inherits this assumption.

2.2 Conditional dependence in latent class models

Violations of local independence are among the oldest concerns in latent class analysis. In diagnostic testing, Vacek (1985) showed that ignoring dependence between tests biases estimated accuracies, and Qu et al. (1996) introduced a latent class model with subject-level random effects, including a class-specific scale, to absorb it. The literature since has developed detection methods—Oberski et al. (2013), in this journal, evaluated the bivariate residual against a chi-square reference together with two methods new to latent class analysis, its parametric-bootstrap version and the score test (modification index), and found that only the latter two perform adequately—and alternative parameterisations: log-linear (fixed pairwise) versus Gaussian random-effects (GRE) formulations, with GRE more parsimonious when dependence is pervasive but numerically more delicate (Xu et al. 2019). Crossed subject-by-rater random effects have been used for ordinal ratings without a gold standard, with robustness checks against misspecified random-effect distributions (Xie et al. 2013). Oberski (2016) distinguished substantive from nuisance dependence, modelling the former as additional discrete latent variables and absorbing the latter in extra parameters, with an application to misclassification of self-reported voting. In model-based clustering of categorical data, Marbac et al. (2016) relaxed local independence by grouping variables into conditionally independent blocks with parsimonious intra-block dependence. Software is available (randomLCA; Beath 2017).

None of this work treats the rater as a factorial object. The rater is a single unit—a test, a pathologist, a survey item—and dependence is either estimated freely or absorbed into a single subject-level factor. In an LLM ensemble the rater has a known design, model $\times$ prompt $\times$ run, and the design tells us where dependence can arise. This is the gap we address.

2.3 LLM annotation: variability, measurement error, and remedies

Three strands of recent work bear on the problem. The first documents instability: accuracy varies substantially across meaning-preserving prompt paraphrases (Liu 2026), and so do the labels themselves (Barrie et al. 2024); accuracy also varies across decoding runs (Atil et al. 2025) and across time as models are updated (Barrie et al. 2025). Large-scale audits of LLM judges find that agreement with human annotation varies widely across tasks and is often limited (Bavaresco et al. 2025), and that high run-to-run reliability can coexist with systematic bias (Norman et al. 2026). The second traces consequences downstream: Baumann et al. (2025) quantified how configuration choices propagate into Type I, II, S and M errors in 2,361 hypothesis tests; Egami et al. (2024) show that annotation error correlated with covariates yields

inconsistent regression estimates regardless of average accuracy, and [Camuffo et al. \(2026\)](#) stress the same point; [Messing \(2026\)](#) extends the argument to evaluation and benchmarking.

The third strand proposes statistical remedies, and two are closest to ours. [Zhang and Martinez \(2025\)](#) reframe LLM stochasticity as measurement error and propose a Bayesian generalised-linear version of Dawid–Skene with covariates on the latent state; their focus is repeated draws from a single pipeline, and raters remain conditionally independent. [Chen et al. \(2026\)](#) take the opposite route: because LLMs share corpora and alignment, their errors may be arbitrarily dependent, so the authors abandon point identification and derive bounds on prevalence, tightened by reporter-specific calibration. Prediction-powered inference ([Angelopoulos et al. 2023](#)) offers a design-based alternative that requires a labelled subsample. Protocol papers ([Camuffo et al. 2026](#)) recommend prompt ensembles, replication and noise-aware aggregation but supply no model in which the dependence between configurations is estimated; diagnostic work such as [Xu et al. \(2026\)](#) decomposes annotation error into task-inherent and model-specific components, again without modelling dependence.

2.4 Positioning

The literature therefore offers independence-based aggregators known to be misspecified for LLM ensembles; bounds that discard the information contained in the ensemble’s design; and protocols without an accompanying inferential model. CRE-LCM occupies the middle ground. It gives dependence the structure implied by how the annotations were generated, which yields point estimates under stated conditions (with identification supported empirically by simulation), interpretable variance components that translate directly into annotation-design decisions, and a natural home for the covariate and multiclass extensions already present in the existing approaches. Methodologically it sits in the ADAC tradition of detecting and relaxing local independence in latent class and mixture models ([Oberski et al. 2013](#); [Marbac et al. 2016](#); [Oberski 2016](#)) and of modelling annotator heterogeneity as a statistical object ([Salter-Townshend and Murphy 2014](#)).

3 Model

3.1 Setting and notation

Items $i = 1, \dots, N$ have an unobserved binary class $c_i \in \{0, 1\}$ (K classes in [Section 3.5](#)) with prevalence $\pi_1 = \Pr(c_i = 1)$. Each item is annotated by every configuration (m, p, r) of M models, P prompts and R decoding runs, giving binary labels Y_{impr} . Runs are exchangeable within a model–prompt pair, so the sufficient statistic is the count $S_{imp} = \sum_r Y_{impr} \in \{0, \dots, R\}$.

3.2 The crossed random-effects latent class model

Conditional on the class and on three item-level random effects, the labels are independent Bernoulli variables with

$$\text{logit Pr}(Y_{impr} = 1 \mid c_i = k, \theta_i, \phi_{im}, \psi_{ip}) = \mu_k + a_{mk} + b_{pk} + \theta_i + \phi_{im} + \psi_{ip}, \quad (1)$$

where $\mu_0 < \mu_1$ are ordered class intercepts, a_{mk} are fixed model effects with the reference model $a_{1k} = 0$, $b_{pk} \sim N(0, \tau_k^2)$ are random prompt effects, and

$$\theta_i \sim N(0, \sigma_\theta^2), \quad \phi_{im} \sim N(0, \sigma_\phi^2), \quad \psi_{ip} \sim N(0, \sigma_\psi^2) \quad (2)$$

are mutually independent: θ_i is the ambiguity of item i shared by all configurations, ϕ_{im} the item $\times$ model effect and ψ_{ip} the item $\times$ prompt effect. Runs are conditionally independent given these effects, so $S_{imp} \mid c_i = k, \theta_i, \phi_{im}, \psi_{ip} \sim \text{Binomial}(R, \text{expit}(\eta_{impk}))$ with η_{impk} the right-hand side of (1), and the marginal likelihood of item i is

$$L_i = \sum_{k=0}^1 \pi_k \int \prod_{m=1}^M \prod_{p=1}^P \binom{R}{S_{imp}} \text{expit}(\eta_{impk})^{S_{imp}} (1 - \text{expit}(\eta_{impk}))^{R - S_{imp}} dF(\theta_i, \phi_i, \psi_i). \quad (3)$$

Nested cases: with $\theta = \phi = \psi = 0$ the model is Dawid–Skene with MP raters whose sensitivities and specificities are $\text{expit}(\mu_1 + a_{m1} + b_{p1})$ and $1 - \text{expit}(\mu_0 + a_{m0} + b_{p0})$; with $\phi = \psi = 0$ it is the one-factor random-effects model of [Qu et al. \(1996\)](#) with a common loading. The random effects play the role of the nuisance-dependence parameters of [Oberski \(2016\)](#), with the difference that their placement is dictated by the design.

3.3 Variance decomposition

On the latent logistic scale the variance of the label propensity of an item over configurations and runs is $\sigma_\theta^2 + \sigma_\phi^2 + \sigma_\psi^2 + \pi^2/3$, the last term being the variance of the standard logistic distribution that generates the run-level Bernoulli outcome. The shares

$$s_{\text{item}} = \frac{\sigma_\theta^2}{V}, \quad s_{\text{model}} = \frac{\sigma_\phi^2}{V}, \quad s_{\text{prompt}} = \frac{\sigma_\psi^2}{V}, \quad s_{\text{run}} = \frac{\pi^2/3}{V}, \quad V = \sigma_\theta^2 + \sigma_\phi^2 + \sigma_\psi^2 + \pi^2/3, \quad (4)$$

have a direct design interpretation: adding models reduces the ϕ share, adding prompts reduces the ψ share, adding runs reduces only the $\pi^2/3$ share, and θ is irreducible by design.

3.4 Identifiability and the merged-class mode

Two questions arise. Are the parameters of CRE-LCM identified without a gold standard, and can they be estimated in practice? We address both with the scenario-B design of Section 4 ($N = 300$, $M = 3$, $P = 5$, $R = 3$; $\sigma_\theta = 1.0$, $\sigma_\phi = 0.8$, $\sigma_\psi = 0.5$), fitted with four chains of 2,000 iterations.

Formal identification. The crossed effects ϕ_{im} and ψ_{ip} are identified by the replication structure: each item is observed under P prompts within every model and under M models within every prompt, so item-by-model and item-by-prompt deviations are separable from the run-level binomial noise whenever $M \geq 2$, $P \geq 2$ and $R \geq 2$. The class intercepts are identified up to label permutation, which we resolve by ordering $\mu_0 < \mu_1$; the model effects are identified relative to a reference model. Without random effects ($\theta = \phi = \psi = 0$) the model is Dawid–Skene with conditionally independent raters, a model that Jones et al. (2010) show to be locally identified with three conditionally independent tests in a single population; in this special case the $R = 3$ runs of each model–prompt pair are also i.i.d. given the class, which meets the condition of Nguyen et al. (2023)—at least $2K - 1$ i.i.d. noisy labels per item for a K -class multinomial mixture—for $K = 2$. For random-effect (latent-trait) extensions, Jones et al. (2010) give only a necessary parameter-counting condition, at least four tests in a single population, which the $MPR = 45$ configurations of our design meet. We give no formal identifiability proof for the full model and rely on the simulation study of Section 4 for empirical support. The delicate parameter is σ_θ : an item that every configuration labels 1 can be explained either as an item of class 1 or as an ambiguous item with a large θ_i , so item ambiguity and class membership are partially confounded. A related substitution between the subject effect and the latent classes appears in Qu et al. (1996), where extreme values of the subject effect absorb the “unequivocal” quasi-classes that a conditional-independence model would need.

The merged-class mode. In practice the more consequential problem was not weak identification but multimodality. With Stan’s default random initialisation, one chain in four converged to a mode in which the mixture has collapsed: prevalence near 0.70, class intercepts almost equal (both between -0.4 and 0), model effects with reversed signs and inflated prompt variances, all heterogeneity being carried by the item effects. The remaining chains agreed with each other and with the truth. Pooling across chains—as a naive analysis would—yields a prevalence of 0.42 [0.27, 0.74] for a true value of 0.30 and an $\hat{R}$ of 1.81; the estimate is an average of two modes. This mode is distinct from label switching (the classes do not exchange; they merge) and is not removed by standard relabelling.

Remedies compared. Table 2 summarises five remedies. Two that seem natural do not work. A class-specific scale for θ_i , in the spirit of Qu et al. (1996), makes matters worse: with $M = 3$ the extra scale trades off against the mixing proportion and the four chains stop at prevalences of 0.17, 0.23, 0.33 and 0.51, the highest-density chain being the one at 0.17. Dropping θ_i altogether—using only the crossed effects—leaves the merged mode in place and pushes the shared item variance into the item-by-model component, whose SD rises from 0.80 to 1.06 [0.97, 1.16] and takes every variance share off target. A tighter prior on σ_θ changes nothing; the data dominate it. Two remedies work. Anchoring 5% of items (15 of 300) with known labels removes the merged mode entirely: all chains agree, all intervals cover the truth, prevalence 0.31 [0.22, 0.38], σ_θ 0.93 [0.77, 1.12], and the posterior labels are the best of any variant (accuracy 0.913, Brier 0.063 over all items; 0.909 and 0.067 on the 285 non-anchored items). Initialising the chains at the Dawid–Skene solution, without changing the model, also removes it: prevalence 0.32 [0.26, 0.39], intercepts $(-2.08, 1.20)$ for $(-2.00, 1.50)$, model

effects recovered, accuracy 0.913, Brier 0.084, in 2.4 minutes against 10 for random starts. The base model is therefore estimable without any gold standard once the sampler starts near the data-consistent mode; we retain the Dawid–Skene initialisation throughout and, in Section 3.6, integrate θ_i out by quadrature; anchoring is reported as the practical remedy for users who hold a small labelled subset—the situation of Section 5.

A residual bias. In every chain that finds the correct mode, whatever the variant, σ_θ is shrunk by 10–20% (0.79–0.92 for a true 1.00; 0.83 [0.67, 1.00] with Dawid–Skene initialisation). Anchoring reduces the shrinkage to 7%. This is the confounding of item ambiguity with class membership noted above and is a property of the likelihood at this design size, not of the estimator; the item share of variance is correspondingly under-estimated (0.14 for 0.19) while the model, prompt and run shares are recovered. We report it as a limitation and examine its dependence on N , P and the strength of dependence in Section 4.

Model comparison. Leave-one-out cross-validation (Vehtari et al. 2017) is not usable here: each item carries $1 + M + P$ random effects of its own, so removing an item changes the posterior substantially ($p_{\text{loo}} \approx 750$ for $N = 300$; 88–93% of Pareto- k values above 0.7 in every variant). Comparisons between variants therefore rest on parameter recovery in simulation, not on predictive criteria.

3.5 Multiclass extension

A K -class version replaces the ordered pair of intercepts by $K - 1$ class-specific linear predictors per configuration, in a multinomial or one-vs-rest formulation, with the same item, item $\times$ model and item $\times$ prompt effects entering each predictor. Identification follows the same replication argument as in Section 3.4, with the class ordering replaced by a reference category. We leave the extension for future work; the binary case is sufficient for the application of Section 5.

3.6 Estimation

We marginalise the discrete class analytically and the shared item effect θ_i numerically, by Gauss–Hermite quadrature with 15 nodes, inside the likelihood of each item; the crossed effects ϕ_{im} and ψ_{ip} are sampled. The continuous parameters are then drawn with the no-U-turn sampler (Stan via cmdstanr; Carpenter et al. 2017) under a non-centred parameterisation, weakly informative priors ($\mu_k \sim N(0, 2)$ ordered; model effects $N(0, 1)$; half-normal(0, 1) on the three SDs; half-normal(0, 0.5) on the prompt-effect SDs τ_k ; Beta(1, 1) on the prevalence) and `adapt_delta` = 0.9, with chains initialised at the Dawid–Skene EM solution. Integrating θ_i out by Gauss–Hermite quadrature is, in effect, the estimation strategy of Qu et al. (1996), who maximised the quadrature-approximated likelihood by EM; we retain it for θ_i while the crossed effects remain in the sampler.

Integrating θ_i out is not a refinement but a necessity. When θ_i is sampled alongside the class, the N item parameters are strongly coupled with the mixing proportion and the class intercepts, and the coupling stiffens with the strength of dependence: in the design of Section 4 the median, over replications, of the smallest bulk effective sample

size among the global parameters (typically that of the prevalence or of σ_θ) falls from 135–290 at $\delta = 0$ to 90–130 at $\delta = 0.5$ and to 5–12 at $\delta \geq 1$, irrespective of N and P , so that chains of 1,500 iterations do not move away from their starting values. Quadrature removes the coupled parameters: on the scenario-B data the bulk effective sample size for the prevalence rises from 13 to about 2,600 per 4,000 post-warm-up draws (four chains of 2,000 iterations), split- $\hat{R}$ is 1.00, the chains agree to three decimals, and recovery improves (prevalence 0.28 [0.22, 0.35] for 0.30; Brier 0.059). Each gradient evaluation costs about five times more, but effective samples per minute increase forty-fold; 10, 15 and 20 nodes agree to within 0.002 on the prevalence and 0.003 on σ_θ , and at $\delta = 0.5$ the quadrature and sampled- θ estimates coincide, so the same estimator is used across the whole simulation study. Two chains of 600 iterations then suffice. The merged-class mode is not entirely removed at the strongest dependence: in the simulation study it recurs in 8% of replications at $\delta = 1.5$, concentrated where the likelihood is sharpest (large N , $P = 5$). We therefore recommend, as a routine diagnostic, running at least two chains and treating a gap above 0.15 between the chain-wise posterior means of the prevalence as a merged-class mode, in which case a small labelled subset should be anchored (Section 3.4). Posterior class probabilities $\Pr(c_i = 1 \mid \text{data})$ and the variance shares are computed in generated quantities; convergence is summarised by split- $\hat{R}$ and bulk and tail effective sample sizes (Vehtari et al. 2021). The Stan programs are listed in Appendix A.

3.7 Computational environment

All computations were run on a single Apple Mac mini (M4 Pro, 64 GB unified memory) hosting a Linux arm64 virtual machine (Ubuntu 26.04.1 LTS, 30 GB RAM, 10 cores allocated). Statistical analyses used R 4.5.3 with cmdstanr 0.9.0 and CmdStan 2.40.0, posterior 1.7.0 and ggplot2 4.0.3; all package versions are pinned in an renv lockfile in the repository. The simulation grid ran on nine parallel workers, one chain per core: about 24 hours for the 2,500 replications with sampled item effects and 36 hours for the 1,300 replications with quadrature; per-replication runtimes are given in Table 5. LLM queries used the Anthropic Messages API (claude-haiku-4-5-20251001), the Gemini API (gemini-3.5-flash, thinking disabled) and Ollama 0.34.4 serving qwen2.5:7b-instruct (Q4_K_M, digest 845dbda0ea48) on the virtual machine; all calls were made between 25 and 26 September 2026 with temperature 0.7, a 16-token output budget and the first digit of the response parsed as the label, plus one temperature-0 pass of prompt 1. Code development was assisted by an AI coding agent, as declared on the title page.

4 Simulation study

4.1 Design

Data are generated from (1)–(2) with $M = 3$ models, $R = 3$ runs, $\pi_1 = 0.30$, $\mu = (-2, 1.5)$, model effects $a_{.0} = (0, 0.3, 0.6)$ and $a_{.1} = (0, -0.1, -0.5)$ (so that later models have both a higher false-positive rate and a lower sensitivity), prompt effects with $\tau_0 = \tau_1 = 0.4$, and random-effect standard deviations $(\sigma_\theta, \sigma_\phi, \sigma_\psi) = \delta \times (1.0, 0.8, 0.5)$.

The dependence level $\delta \in \{0, 0.5, 1, 1.5\}$ is crossed with $N \in \{150, 300, 600\}$ and $P \in \{3, 5\}$, giving 24 cells with 100 replications each. A misspecification arm at $\delta = 1$, $N = 300$, $P = 5$ replaces the three Gaussian random effects by scaled Student- t_4 variables with the same standard deviations. Every replication has a recorded seed (data and sampler) and is stored individually, so the study is reproducible and resumable from the public code repository.

Three estimators are compared: majority vote over the *MPR* labels; the binary Dawid–Skene model with *MPR* raters fitted by EM (Dawid and Skene 1979); and CRE-LCM as described in Section 3.6, initialised at the Dawid–Skene solution. At $\delta \leq 0.5$ the item effects are sampled (two chains of 1500 iterations, 750 warm-up); at $\delta \geq 1$, where the sampled chains do not mix (see below), the shared item effect is integrated out by 15-node Gauss–Hermite quadrature (two chains of 600 iterations, 300 warm-up). The two routes target the same posterior: on 20 replications at $\delta = 0.5$, $N = 300$ fitted both ways, the prevalence estimates differ by -0.001 on average (SD 0.002, correlation 0.995) and every other quantity agrees within Monte Carlo error, and 10, 15 and 20 nodes give identical results on the scenario-B data. For each replication we record the bias of the estimated prevalence, of mean sensitivity and of mean specificity (averaged over the $M \times P$ configurations; for CRE-LCM these are the posterior means of the marginal accuracies implied by the model, integrating over the random effects), classification accuracy and Brier score of the posterior class probabilities against the simulated classes, coverage of the CRE-LCM 95% credible intervals, and the estimated variance shares (4). Replications in which any global parameter of CRE-LCM has split- $\hat{R} > 1.05$ are flagged and their fraction reported per cell.

4.2 Dawid–Skene under dependence

Table 1 gives a first view on a single data set per dependence level. Under conditional independence ($\delta = 0$) Dawid–Skene recovers sensitivity, specificity and prevalence essentially exactly. As δ grows the true marginal accuracies fall, because the item effects push labels away from the class-specific propensities, but the Dawid–Skene estimates do not: at $\delta = 1.5$ the model reports a specificity of 0.82 against a true 0.74 and a prevalence of 0.36 against 0.29. The bias is monotone in δ and is larger for prevalence and specificity than for sensitivity, which the results over the full grid below confirm.

4.3 Identifiability variants

Table 2 reports the five remedies of Section 3.4, fitted to one scenario at $\delta = 1$, $N = 300$, $P = 5$ with four chains of 2000 iterations. With default random initialisation, one of the four chains of the base model settles in the merged-class mode (per-chain prevalence near 0.70, class intercepts almost equal, both between -0.4 and 0 , model effects sign-flipped); the pooled posterior is a mixture of the two modes and $\hat{R}$ on π_1 is 1.81. The class-specific scale (a) makes matters worse: each chain ends at a different prevalence (0.17 to 0.51) and the chain with the highest posterior density is the one furthest from the truth, so the extra loading is not identified separately

Table 1 Dawid–Skene under crossed dependence: one simulated data set per dependence level ($N = 300$, $M = 3$, $P = 5$, $R = 3$, seed 2026). Dependence level δ scales the three random-effect standard deviations $(\sigma_\theta, \sigma_\phi, \sigma_\psi) = \delta \times (1.0, 0.8, 0.5)$. Sensitivity and specificity are averaged over the $M \times P$ configurations

δ	Sensitivity		Specificity		Prevalence		
	true	DS	true	DS	true	DS	MV
0.0	0.710	0.710	0.853	0.856	0.290	0.290	0.290
0.5	0.696	0.720	0.832	0.832	0.290	0.287	0.280
1.0	0.666	0.703	0.787	0.822	0.290	0.327	0.287
1.5	0.638	0.714	0.741	0.819	0.290	0.359	0.307

from the mixing proportion in this design. Omitting θ_i (b) does not remove the mode and biases the remaining components, the item $\times$ model standard deviation rising to 1.06 against 0.80, because the shared item variance is pushed into the crossed effects. The informative prior (d) changes nothing, the data dominating the prior. Anchoring 5% of the items (c) removes the mode entirely, with every interval covering the truth and the best label quality of all variants. Initialising the chains at the Dawid–Skene solution (e) also removes it: all chains agree, the intervals cover the truth, and label quality is close to that of anchoring. Two limitations remain under (e): mixing on π_1 and σ_θ is slow, and σ_θ is shrunk downwards by 10–20% in every chain that sits in the correct mode of every variant, anchoring reducing the shrinkage. We therefore retain the Dawid–Skene initialisation throughout, integrate θ_i out by quadrature in the estimator of Section 3.6, and report anchoring as the practical remedy when a small labelled subset exists.

Table 2 Identifiability variants fitted to one scenario-B data set ($N = 300$, $\sigma_\theta = 1$, $\sigma_\phi = 0.8$, $\sigma_\psi = 0.5$, true $\pi_1 = 0.30$, $\mu = (-2, 1.5)$; four chains of 2000 iterations). Posterior means, with 95% credible intervals for π_1 and σ_θ ; $\hat{R}$ is the largest split- $\hat{R}$ over the global parameters

Variant	π_1	μ_0, μ_1	σ_θ	Acc.	Brier	max $\hat{R}$
Majority vote	–	–	–	0.890	0.110	–
Dawid–Skene	–	–	–	0.887	0.109	–
Base	0.42 [0.27, 0.74]	-1.84, 0.65	0.86 [0.69, 1.09]	0.897	0.121	1.81
(a) class-specific σ_θ	0.31 [0.14, 0.60]	-1.52, 1.15	1.17 / 0.80	0.897	0.107	2.67
(b) θ omitted	0.41 [0.28, 0.71]	-1.62, 0.92	0.00 [0.00, 0.00]	0.897	0.114	1.61
(c) 5% anchored	0.31 [0.22, 0.38]	-2.02, 1.23	0.93 [0.77, 1.12]	0.913	0.063	1.44
(d) HN(0, 0.5) prior	0.42 [0.27, 0.75]	-1.65, 0.82	0.86 [0.70, 1.03]	0.883	0.123	1.64
(e) DS initialisation	0.32 [0.26, 0.39]	-2.08, 1.20	0.83 [0.67, 1.00]	0.913	0.084	1.20

In (a) the two entries are $\sigma_{\theta,0} / \sigma_{\theta,1}$. Accuracy and Brier score of the posterior class probabilities against the simulated true classes (all 300 items; for (c) the 15 anchored items are included).

4.4 Results over the grid

Figure 1 shows the bias of prevalence, sensitivity and specificity; Figure 2 the accuracy and Brier score of the posterior labels; Figure 3 the coverage of the CRE-LCM 95% intervals; and Figure 4 the recovery of the variance shares at $N = 300$, $P = 5$. Tables 3–5 give the per-cell values.

Dawid–Skene. The single-data-set pattern of Table 1 holds over the grid. Under conditional independence Dawid–Skene is unbiased for every quantity at every N and P . Its prevalence bias is about +0.003 at $\delta = 0.5$, +0.05 at $\delta = 1$ and +0.10 to +0.11 at $\delta = 1.5$, and its specificity bias +0.04 and +0.09 at $\delta = 1$ and 1.5, essentially independent of N and P : the bias is a property of the misspecification, not of the sample size. Sensitivity is much less affected (at most +0.03). Majority vote is less biased for prevalence (at most +0.05) but has no accuracy estimates.

CRE-LCM at $\delta \leq 0.5$. Under conditional independence and weak dependence CRE-LCM is unbiased for every quantity: prevalence bias at most 0.004, sensitivity and specificity bias below 0.004, coverage of the prevalence interval 1.00, and Brier scores below Dawid–Skene’s in every cell. Convergence is unproblematic (at most 20% of replications flagged at $\delta = 0.5$, median bulk ESS above 90). A prevalence coverage of 1.00 indicates that the intervals are conservative at weak dependence.

CRE-LCM at $\delta \geq 1$. With the shared item effect integrated out, the model removes most of the Dawid–Skene bias. At $\delta = 1$ the prevalence bias is +0.006 to +0.018 (Dawid–Skene: +0.05), the specificity bias -0.004 to -0.015 (Dawid–Skene: +0.04), and the coverage of the prevalence interval 0.96 to 1.00. At $\delta = 1.5$ the prevalence bias is +0.027 to +0.065 against +0.10 to +0.11 for Dawid–Skene, with coverage 0.94 to 0.99; the largest values occur at $N = 150$. The one quantity that CRE-LCM does not recover well is mean sensitivity, which it underestimates by 0.01 to 0.05 at $N \geq 300$ and by 0.04 to 0.09 at $N = 150$. This is a small-sample bias specific to the minority class: with a prevalence of 0.30 the class-1 intercept and effects are estimated from few items, the ordered constraint and the prior shrink μ_1 towards zero, and the marginal sensitivity falls. It decreases with N (at $\delta = 1.5$, from -0.08 at $N = 150$ to -0.04 at $N = 600$ once merged-mode replications are excluded; the exception is the cell $N = 600$, $P = 5$, discussed below), coverage of the sensitivity interval stays at 0.89 to 0.98, and it is conservative and smaller than the Dawid–Skene overstatement (+0.04 to +0.09 of specificity, +0.03 of sensitivity) in the opposite direction. The posterior labels are markedly better calibrated: the Brier score is 0.05 to 0.07 against 0.10 for Dawid–Skene at $\delta = 1$ and 0.12 to 0.14 against 0.20 at $\delta = 1.5$ (a ratio of about 0.65 in every cell), while accuracy is higher by one to four points. The variance shares are recovered: the model and prompt shares with bias below 0.02 and coverage 0.87 to 0.97, the item share with a positive bias of 0.01 to 0.05 that shrinks with N (Fig. 4), the opposite sign to the shrinkage of σ_θ seen when θ_i is sampled (Section 3.4). Convergence at $\delta = 1$ is satisfactory (5% to 20% of replications flagged, median smallest bulk ESS 89 to 162 from 600 draws); at $\delta = 1.5$ mixing remains slower (36% to 77% flagged, median smallest ESS 18 to 64) and the merged-class mode reappears in 8% of replications, concentrated at $N = 600$, $P = 5$ (18 of 100), that is, where the likelihood is sharpest. We report both figures: over all replications the $\delta = 1.5$ prevalence bias is +0.043; excluding the merged-mode replications, identified by a gap above 0.15

between the chain-wise posterior means of the prevalence, it is $+0.032$. In the cell $N = 600$, $P = 5$ the 18 merged-mode replications account for the whole difference between the two estimators: over all replications the CRE-LCM prevalence RMSE (0.112) exceeds Dawid-Skene's (0.108) and the sensitivity bias is -0.065 ; excluding them, the prevalence bias is $+0.016$ with RMSE 0.083 (Dawid-Skene $+0.102$, 0.106), the sensitivity bias -0.036 , the specificity bias -0.019 , and coverage 0.93, 0.88 and 0.89 for prevalence, sensitivity and specificity (notes to Tables 3 and 4). The diagnostic of Section 3.6, followed by anchoring a small labelled subset as in Section 3.4, is the remedy at this strength of dependence. Runtime per replication with quadrature is 3.5 and 6 minutes at $N = 150$ ($P = 3$ and 5), 8 and 14 minutes at $N = 300$, and 21 and 37 minutes at $N = 600$ on one core.




Fig. 1 Bias of the estimated prevalence (top), mean sensitivity (middle) and mean specificity (bottom) by dependence level δ for majority vote, Dawid-Skene and CRE-LCM; columns are sample sizes, line types are the number of prompts. Means over 100 replications per cell




Fig. 2 Classification accuracy (top) and Brier score (bottom) of the posterior class probabilities by dependence level δ ; layout as in Fig. 1




Fig. 3 Coverage of the CRE-LCM 95% credible intervals for prevalence, mean sensitivity, mean specificity and the three variance shares, by dependence level δ and sample size; the dotted line marks the nominal level. Coverage of a share is undefined at $\delta = 0$, where its true value is zero




Fig. 4 Posterior means of the four variance shares over 100 replications at $N = 300$, $P = 5$, by dependence level δ ; crosses mark the true shares

Table 3 Prevalence estimation over the simulation grid ($M = 3$, $R = 3$; n replications per cell): bias and root mean squared error of majority vote (MV), Dawid-Skene (DS) and CRE-LCM, and coverage of the CRE-LCM 95% credible interval

P	N	δ	n	MV		DS		CRE-LCM		
				bias	RMSE	bias	RMSE	bias	RMSE	cov.
3	150	0.0	100	-0.001	0.002	0.000	0.001	0.003	0.003	1.00
		0.5	100	-0.005	0.018	0.004	0.011	0.002	0.010	1.00
		1.0	100	0.016	0.043	0.053	0.072	0.018	0.063	1.00
		1.5	100	0.045	0.065	0.103	0.120	0.065	0.110	1.00
	300	0.0	100	0.000	0.002	0.000	0.001	0.001	0.002	1.00
		0.5	100	-0.003	0.014	0.003	0.009	0.001	0.007	1.00
		1.0	100	0.013	0.032	0.047	0.059	0.006	0.053	0.96
		1.5	100	0.053	0.064	0.107	0.116	0.038	0.089	0.99
	600	0.0	100	-0.001	0.002	0.000	0.001	0.001	0.001	1.00
		0.5	100	-0.005	0.012	0.003	0.008	0.000	0.004	1.00
		1.0	100	0.018	0.035	0.052	0.061	0.006	0.052	0.98
		1.5	100	0.047	0.058	0.106	0.112	0.026	0.086	0.99
5	150	0.0	100	0.000	0.000	0.000	0.000	0.003	0.003	1.00
		0.5	100	-0.002	0.009	0.004	0.009	0.004	0.008	1.00
		1.0	100	0.013	0.034	0.055	0.066	0.009	0.050	1.00
		1.5	100	0.044	0.062	0.104	0.117	0.049	0.103	0.99
	300	0.0	100	0.000	0.000	0.000	0.000	0.001	0.001	1.00
		0.5	100	-0.004	0.009	0.002	0.006	0.001	0.006	1.00
		1.0	100	0.013	0.029	0.054	0.063	0.001	0.028	0.99
		1.5	100	0.045	0.053	0.107	0.115	0.030	0.094	0.94
	600	0.0	100	0.000	0.000	0.000	0.000	0.001	0.001	1.00
		0.5	100	-0.005	0.009	0.001	0.004	0.001	0.004	1.00
		1.0	100	0.014	0.025	0.049	0.054	0.018	0.071	0.99
		1.5 ^a	100	0.044	0.050	0.104	0.108	0.048	0.112	0.94

^a 18 of 100 replications in this cell sit in the merged-class mode (chain-wise posterior prevalence differing by more than 0.15); excluding them, the CRE-LCM prevalence bias is 0.016, the RMSE 0.083 and the coverage 0.93.

Table 4 Bias of mean sensitivity and mean specificity (averaged over the $M \times P$ configurations) for Dawid–Skene (DS) and CRE-LCM, with coverage of the CRE-LCM 95% intervals

P	N	δ	Sensitivity			Specificity		
			DS	CRE	cov.	DS	CRE	cov.
3	150	0.0	-0.001	-0.002	0.93	0.000	0.001	0.96
		0.5	0.000	-0.004	0.89	0.002	-0.002	0.93
		1.0	-0.002	-0.043	0.89	0.041	-0.015	0.94
		1.5	0.026	-0.088	0.96	0.090	-0.028	0.95
	300	0.0	-0.001	-0.002	0.93	0.001	0.001	0.94
		0.5	0.002	-0.002	0.90	0.002	-0.001	0.90
		1.0	0.003	-0.023	0.91	0.035	-0.011	0.87
		1.5	0.027	-0.053	0.98	0.093	-0.017	0.97
	600	0.0	0.001	0.000	0.97	0.000	0.000	0.93
		0.5	0.000	-0.003	0.94	0.002	-0.001	0.85
		1.0	-0.002	-0.020	0.93	0.038	-0.015	0.91
		1.5	0.022	-0.048	0.93	0.091	-0.021	0.95
5	150	0.0	0.000	-0.001	0.94	0.000	0.000	0.94
		0.5	0.000	-0.004	0.95	0.002	0.000	0.93
		1.0	-0.011	-0.035	0.92	0.038	-0.015	0.93
		1.5	0.019	-0.078	0.96	0.085	-0.026	0.96
	300	0.0	-0.002	-0.002	0.95	0.000	0.000	0.96
		0.5	0.001	-0.002	0.90	0.003	0.001	0.95
		1.0	-0.006	-0.012	0.95	0.038	-0.004	0.93
		1.5	0.019	-0.050	0.93	0.086	-0.021	0.89
	600	0.0	-0.001	-0.001	0.96	0.000	0.000	0.96
		0.5	0.000	-0.004	0.88	0.001	0.000	0.90
		1.0	-0.003	-0.020	0.91	0.034	-0.009	0.95
		1.5 ^a	0.019	-0.064	0.90	0.087	-0.047	0.87

^a Excluding the 18 merged-mode replications (see Table 3), the CRE-LCM sensitivity bias is -0.035 with coverage 0.88 and the specificity bias -0.019 with coverage 0.89.

4.5 Misspecification

Table 6 compares the Gaussian and t_4 data-generating processes at $\delta = 1$, $N = 300$, $P = 5$. Heavy-tailed random effects do not degrade CRE-LCM: prevalence bias +0.009 against +0.001 under the Gaussian process (Dawid–Skene: +0.037 and +0.054), accuracy 0.94 against 0.93, Brier 0.046 against 0.052, coverage of the prevalence interval 1.00 against 0.99, and 8% of replications flagged against 10%. The one visible effect of the misspecification is on the item share, which is underestimated by 0.036 with coverage 0.67 (Gaussian: +0.012, 0.97): the Gaussian working model attributes part of the heavy-tailed item variance to the crossed effects and to the run-level residual. Point estimates of prevalence and accuracy, and the posterior labels, are therefore robust to the tail behaviour of the random effects at this design size; the variance decomposition is the quantity to interpret with care when heavy tails are suspected.

Table 5 Posterior label quality: classification accuracy and Brier score of the item-level class probabilities (MV and DS: EM posterior; CRE-LCM: posterior mean), the fraction of replications with any split- $\hat{R} > 1.05$ among the CRE-LCM global parameters, and mean CRE-LCM runtime per replication (minutes on one core; $\delta \leq 0.5$: sampled item effects, two chains of 1500 iterations; $\delta \geq 1$: item effect integrated by 15-node Gauss–Hermite quadrature, two chains of 600 iterations)

P	N	δ	Accuracy			Brier			$\hat{R} > 1.05$	min
			MV	DS	CRE	MV	DS	CRE		
3	150	0.0	0.999	1.000	1.000	0.001	0.000	0.000	0.00	1.1
		0.5	0.981	0.985	0.987	0.019	0.014	0.010	0.05	1.1
		1.0	0.900	0.893	0.908	0.100	0.103	0.072	0.20	3.0
		1.5	0.808	0.792	0.808	0.192	0.203	0.144	0.45	4.0
	300	0.0	0.999	1.000	1.000	0.001	0.000	0.000	0.01	2.7
		0.5	0.982	0.987	0.988	0.018	0.012	0.009	0.13	2.7
		1.0	0.902	0.895	0.909	0.098	0.100	0.066	0.06	7.3
		1.5	0.812	0.795	0.821	0.188	0.199	0.129	0.42	9.4
	600	0.0	0.999	1.000	1.000	0.001	0.000	0.000	0.02	7.5
		0.5	0.982	0.986	0.989	0.018	0.012	0.009	0.17	7.5
		1.0	0.899	0.894	0.908	0.101	0.099	0.071	0.05	20.3
		1.5	0.814	0.796	0.826	0.186	0.195	0.130	0.36	21.9
	5	150	0.0	1.000	1.000	0.000	0.000	0.000	0.00	1.8
			0.5	0.989	0.992	0.011	0.008	0.006	0.18	1.8
			1.0	0.907	0.895	0.093	0.103	0.060	0.14	4.8
			1.5	0.819	0.801	0.181	0.197	0.129	0.77	7.3
		300	0.0	1.000	1.000	0.000	0.000	0.000	0.01	4.4
			0.5	0.988	0.991	0.012	0.008	0.006	0.19	4.6
			1.0	0.906	0.897	0.094	0.100	0.052	0.10	11.5
			1.5	0.813	0.794	0.187	0.202	0.123	0.56	17.2
	600	0.0	1.000	1.000	1.000	0.000	0.000	0.000	0.03	11.9
		0.5	0.988	0.992	0.992	0.012	0.008	0.007	0.20	12.2
		1.0	0.908	0.902	0.921	0.092	0.095	0.061	0.05	32.5
		1.5	0.820	0.802	0.821	0.180	0.193	0.142	0.61	41.1

5 Application: policy-topic coding of US congressional bills

5.1 Data

We apply the model to the classification of US congressional bill titles by policy topic. The Congressional Bills Project (Wilkerson et al. 2025) provides the title of every bill introduced in the 80th–114th Congresses (1947–2016) together with a major-topic code assigned by trained human coders under the Comparative Agendas Project (CAP) codebook (Jones et al. 2025), which makes the human code a genuine reference label rather than a label produced by the raters under study. From the 37,007 bills of the 111th to 114th Congresses (2009–2016), the most recent covered by the dataset, we drew a stratified sample of 800 titles: 240 coded as *Health* (CAP major topic 3) and 560 from the other nineteen substantive topics in proportion to their frequency, giving a prevalence of 0.30 for the binary target “the bill’s main subject is health”. Health accounts for 11.6% of the eligible bills; the stratification raises the minority class to the

Table 6 Recovery of the variance shares by CRE-LCM at $N = 300$, $P = 5$ (bias, RMSE and 95% coverage), with the misspecification arm in which the three item random effects follow a scaled t_4 distribution at $\delta = 1$; the last three columns give the CRE-LCM prevalence bias, accuracy and Brier score in the same cells

DGP	δ	n	Item share			Model share			Prompt share			CRE-LCM			
			bias	RMSE	cov.	bias	RMSE	cov.	bias	RMSE	cov.	prev.	bias	acc.	Brier
Gaussian	0.0	100	0.002	0.003	—	0.003	0.004	—	0.005	0.006	—	0.001	1.000	0.000	
Gaussian	0.5	100	0.000	0.011	0.94	0.001	0.010	0.93	-0.001	0.008	0.95	0.001	0.993	0.006	
Gaussian	1.0	100	0.012	0.031	0.97	0.000	0.014	0.96	-0.001	0.009	0.94	0.001	0.930	0.052	
Gaussian	1.5	100	0.032	0.063	0.87	-0.011	0.024	0.91	-0.004	0.010	0.92	0.030	0.833	0.123	
t_4	1.0	100	-0.036	0.046	0.67	-0.004	0.015	0.90	0.001	0.010	0.91	0.009	0.941	0.046	

prevalence of the simulation study and keeps the number of positive items adequate for the class-specific parameters (Section 4.4). Titles are short (median 25 words), which keeps the cost of a full factorial annotation low. Two hundred items were set aside at random as a validation set whose human codes are never used in estimation and which is never anchored; the remaining 600 form the working set (174 Health, 0.29); all fits use the LLM labels of all 800 bills, the primary fit uses no human code, and the anchored variant uses the codes of 30 working bills only. The sample, the codes and all annotations are archived (see the data availability statement).

5.2 Annotation design

Each title was labelled by every configuration of $M = 3$ model families, $P = 5$ prompts and $R = 3$ decoding runs at temperature 0.7, that is, 45 labels per item and 36,000 labels in total, plus one pass at temperature 0 with the first prompt for each model. The five prompts are semantically equivalent formulations of the same binary question and differ in framing: a bare question (P1), a codebook definition of the Health topic (P2), an explicit contrast with the other policy areas (P3), an analyst role (P4), and an instruction to consider the bill’s purpose before answering (P5). All ask for a single digit; the first digit of the response is taken as the label and a response without a digit is retried once and then recorded as missing. The templates are given verbatim in Online Resource 1; they were written by the author rather than generated by a model. Every call used a 16-token output budget, and every raw response is stored with its timestamp, model version, prompt, run and temperature.

The three model families were chosen to span the design space that the method is meant to describe rather than to maximise accuracy. First, each family comes from a different developer with a different training corpus and alignment procedure (Anthropic, Google, Alibaba), so that the item-by-model effect ϕ_{im} captures genuine between-family dependence rather than variation among versions of one model. Second, the two commercial models are the efficiency tier of their families (Claude Haiku 4.5, Gemini 3.5 Flash), which is the tier practitioners use for annotation at scale: [Carlson and Burbano \(2026\)](#) found that an efficiency-oriented model (GPT-4o) outperformed the larger models they tested on binary classification at a fraction of their cost, although the smallest models they tested were markedly less accurate; cost is what makes the model $\times$ prompt $\times$ run design affordable. Third, one family is an open-weights model run locally (Qwen 2.5 7B Instruct, 4-bit quantised, served by

Ollama): its weights and digest are archived, so this arm of the ensemble is exactly reproducible, which no API-served model can guarantee once its version is retired (Barrie et al. 2025). Larger or reasoning-oriented models were deliberately excluded: they raise cost by one to two orders of magnitude, their suitability for plain classification is unclear (Carlson and Burbano 2026), and the question of the paper is not which model is best but how dependent the members of an ensemble are. Model identifiers, quantisation, decoding parameters, query dates, call counts, missing rates and cost are given in Table 7; all raw responses are archived.

Table 7 reports, per provider, the model identifier, quantisation, decoding settings, query dates, number of calls, missing rate and cost. The three providers answered 12,800 calls each; Anthropic and the local model returned a digit in every case, Gemini in all but 10 (0.08%). Total API cost was USD 3.31; the local model cost nothing. The ten missing labels are treated as missing at random, which the binomial formulation accommodates through a reduced run count R_{imp} for the affected configuration–item pairs.

Table 7 LLM configurations used in the application. Calls count the $800 \times 5 \times 3$ labels at temperature 0.7 plus the 800-call temperature-0 pass of prompt 1; missing = answers without a 0/1 digit after one retry; cost at list prices from the logged token counts (Haiku 4.5: USD 1.00 / 5.00 per million input / output tokens; Gemini 3.5 Flash: USD 1.50 / 9.00)

Provider	Anthropic	Google	Alibaba (open weights)
Model identifier	claude-haiku-4-5-20251001	gemini-3.5-flash	qwen2.5:7b-instruct
Version / digest	as served (API string)	as served, thinking off	digest 845dbda0ea48
Quantisation	–	–	Q4_K_M, 7.6B
Endpoint	Messages API	Gemini API	Ollama 0.34.4 (local)
Temperature	0.7 (and 0)	0.7 (and 0)	0.7 (and 0)
Max output tokens	16	16	16
Query dates	25 Sept 2026	25–26 Sept 2026	25 Sept 2026
Calls	12,800	12,800	12,800
Missing (%)	0.00	0.08	0.00
Cost (USD)	1.63	1.68	0

5.3 Results

Agreement. Fleiss’ κ across the fifteen prompt $\times$ run configurations is 0.85 for Claude Haiku, 0.97 for Gemini Flash and 0.84 for Qwen, and 0.84 over all 45 configurations; mean pairwise agreement is 0.94, 0.99 and 0.94 within the three models and 0.92 to 0.94 between them (Table 8). Agreement is high: the configurations rarely disagree on a title, and when they do, they disagree on the same titles—the boundary cases the design was meant to expose. The 116 bills (14.5%) that receive between 6 and 39 Health labels out of 45 are mostly Medicare and Medicaid financing, veterans’ health and drug-control bills: 72% of them are coded Health by CAP and the rest mainly law and crime, defence or social welfare.

Dawid–Skene. Treating the 45 configurations as independent raters, Dawid–Skene estimates a prevalence of 0.331 for a sample value of 0.300, mean sensitivity 0.866 and mean specificity 0.994; its posterior class probabilities average 0.321 on the working

set (CAP 0.291) and 0.370 on the validation set (CAP 0.330). The +0.03 bias in prevalence lies between the values the simulation study produces at $\delta = 0.5$ and $\delta = 1$ (Table 3), and the pattern is the same: the raters’ near-unanimity on easy items is read as near-perfect specificity, and the items on which they jointly err are read as positives.

CRE-LCM. Without any human label, the crossed random-effects model estimates a prevalence of 0.315 [0.284, 0.347], whose interval covers the sample value; the bias is halved, from +0.031 to +0.015. The variance decomposition is the substantive result (Table 9, reported from the anchored fit, which converges; the unanchored point estimates agree within 0.01): the shared item effect accounts for 0.59 [0.53, 0.65] of the variance of the label propensity, the item-by-model effect for 0.19 [0.15, 0.23], the item-by-prompt effect for 0.16 [0.13, 0.19] and decoding randomness for only 0.06 [0.05, 0.07]. Three-fifths of the disagreement in this ensemble reflects properties of the text that all configurations share; a further third reflects which model and which prompt were used, in roughly equal measure; and the repeated decoding runs, which a “self-consistency” protocol would multiply, contribute almost nothing. The class intercepts (−14.5 and +7.1) and the item scale (5.6) are large because the labels are close to deterministic: on this task the logistic scale is stretched. We keep the priors of Section 3.6, chosen for the simulation, as the primary specification and report a refit with wider priors below. In the unanchored fit the item scale and the shares mix slowly on this stretched scale (bulk ESS 8 to 16 from 6,000 draws, $\hat{R}$ up to 1.46), while the prevalence mixes well (ESS 1,160); all four chains sit at the same values and none shows the merged-class mode.

Anchoring. With 5% of the working items (30) anchored to their CAP code, the prevalence estimate is 0.311 [0.281, 0.343] and the shares are unchanged (0.59, 0.19, 0.16 and 0.06 against 0.60, 0.18, 0.16 and 0.06 without anchors); on the 570 non-anchored items the posterior labels reach an accuracy of 0.960 and a Brier score of 0.038. The anchors matter here for mixing rather than for the estimates: every parameter converges ($\hat{R} \leq 1.04$, bulk ESS of the shares 134 to 349).

Prior sensitivity. Refitting both models with wider priors ($N(0, 10)$ on the intercepts, half-normal(0, 5) on the three standard deviations and on τ_k ; four chains of 4,000 iterations) moves the latent scale but not the substance (Table B1, Appendix B). With near-unanimous labels the likelihood is nearly flat in the scale of the latent effects—the data are close to quasi-complete separation—so under the wide priors the class intercepts drift to about −22 and +11, σ_θ to 7.8–8.5, the chains mix poorly ($\hat{R}$ up to 1.9), and the absolute shares follow the prior: the item share rises to 0.64–0.69 and the run share falls to 0.03–0.04. The ranking item $\gg$ model $\approx$ prompt $\gg$ run does not change, nor do the prevalence (0.31–0.32), the accuracy and Brier score against the CAP codes (to within 0.002), the per-configuration sensitivities and specificities (to within 0.025) and the posterior predictive check.

Validation. On the 200 held-out items, majority vote, Dawid–Skene and CRE-LCM classify 97.0%, 96.0% and 97.0% correctly, with Brier scores of 0.030, 0.040 and 0.031; on the 570 non-anchored working items the figures are 95.4%, 95.6% and 96.0% with Brier scores of 0.046, 0.044 and 0.039 (Table 10). CRE-LCM’s gain over Dawid–Skene is in calibration rather than in accuracy, as in the simulation; against majority vote,

whose hard labels make the Brier score equal to the error rate, the gain appears only on the larger working set. The posterior probabilities of both model-based methods are close to 0 or 1 for almost every bill, so the calibration plot (Fig. 5) is dominated by its two end bins. Table 11 compares, for each of the 15 model $\times$ prompt configurations, the sensitivity and specificity implied by CRE-LCM with the values computed against the CAP codes of all 800 bills: the model-based values track the observed ones within 0.03 on average (at most 0.08 for sensitivity and 0.04 for specificity), whereas Dawid–Skene overstates specificity by 0.02 to 0.05 (0.036 on average) and compresses its range across configurations to 0.021, against 0.043 in the CAP values. CRE-LCM narrows but does not remove the gap in specificity (+0.026 on average, range 0.024), and the CAP value falls inside its 95% interval for only 6 of the 15 sensitivities and 2 of the 15 specificities: bills that the ensemble labels Health but CAP codes under another topic are, for any latent class model, positives with an ambiguous item rather than errors of the raters. The posterior predictive check (Fig. 6) compares the observed distribution of the number of Health labels per item across the 45 configurations with its replicated distribution under CRE-LCM and under Dawid–Skene: CRE-LCM reproduces the U-shaped distribution with its thin but non-empty middle (5.6% of bills with 15 to 30 Health labels, replicated 4.9%; 20.4% with 40 or more, replicated 20.6%), whereas Dawid–Skene, whose raters are independent given the class, predicts almost no intermediate counts (0.006% with 15 to 30) and places the Health bills at 36 to 41 labels instead of 42 to 45 (13.8% with 40 or more).

Temperature 0. The temperature-0 pass with prompt 1 agrees with the temperature-0.7 majority for the same prompt on 99.9%, 100% and 98.5% of items for Claude Haiku, Gemini Flash and Qwen, with the same accuracy against the CAP codes (0.94, 0.96 and 0.92); the three models at temperature 0 disagree with one another on 9.6% of items, so a deterministic decoding setting removes the smallest variance component and leaves the two largest untouched.

Shared construct bias. All six validation errors of CRE-LCM are confident (posterior probability of Health above 0.99) and point in the same direction: the ensemble labels as Health bills that CAP coded under another topic—two drug-control bills (controlled-substance analogues, medical marijuana), a mental-health screening of military recruits, a medical-isotope supply bill, a restriction on research involving wards of the State and a fuel-tax refund for mobile mammography vehicles. The per-configuration comparison shows the same pattern, with the CAP values often outside the CRE-LCM intervals. Both reflect a difference between the LLMs’ shared notion of “health” and the CAP coding rule, which assigns such bills to their legal or institutional home, and the difference is common to every configuration. A latent class model cannot identify a bias shared by all raters without a gold standard: the bias becomes part of the latent class, which here is health as the LLMs read it rather than the CAP topic. Anchoring a small coded subset is the remedy, because the anchors tie the latent class to the reference coding; but it works only if the anchors include boundary bills. The 30 anchors drawn at random here contain one non-Health bill with a substantial share of Health labels, and the anchored fit makes the same six validation errors. This parallels the reliability-without-validity finding of Norman et al. (2026), who show that an LLM judge can be almost perfectly consistent from run to run while

remaining systematically biased: within-model run-to-run consistency does not guarantee validity. Here the same holds across models—the LLM raters agree with one another more closely than with the CAP coding—and no model of their agreement alone can recover the difference.

5.4 Design implications for this task

The decomposition answers the budget question of Section 6 directly. For these titles, a fourth decoding run would reduce the residual share of 0.06 and nothing else; a fourth prompt or a fourth model family would each act on a share of roughly 0.17; and the item share of 0.59 can be reduced only by sharpening the coding rule—for example by deciding in advance how veterans’ health or health-related welfare bills are to be coded—or by adjudicating the flagged items by hand. The posterior class probabilities identify few such items: only 5 of the 600 working items have a posterior between 0.2 and 0.8, and none of the six validation errors of CRE-LCM (one of the eight of Dawid-Skene) falls among the flagged items. The errors against the CAP codes are the confident, shared ones described above. The agreement count is the better guide for adjudication: four of the six errors are among the 32 validation bills with 6 to 39 Health labels out of 45, although their Health posterior exceeds 0.8.

Table 8 Agreement among the LLM configurations on the 800 bills: Fleiss’ κ and mean pairwise agreement within each model (15 configurations: 5 prompts $\times$ 3 runs) and over all 45 configurations; percentage of labels equal to 1 and of missing (unparsable) answers

Model	Configs	Fleiss’ κ	Agreement	% label 1	% missing
anthropic	15	0.848	0.938	28.6	0.00
gemini	15	0.972	0.988	32.8	0.08
ollama	15	0.844	0.940	25.9	0.00
all 45 configurations	45	0.842	0.935	29.1	0.03

6 Discussion

What the variance decomposition is for. The three random-effect shares answer a question that agreement statistics cannot: where should the next unit of annotation budget go? Decoding runs only reduce the residual share; they cannot touch item, model or prompt variance. When the item-by-model share dominates, a fourth model family is worth more than any number of re-runs; when the item-by-prompt share dominates, the prompt family should be widened. The shared item share is the ceiling: it is the part of disagreement that reflects genuine ambiguity in the text and is reduced only by better coding rules or human adjudication. We therefore recommend that studies using LLM annotation report the three shares with their intervals next to the customary κ , and that “temperature 0, single prompt, single model” be treated as a design with no replication rather than as a deterministic instrument.




Fig. 5 Reliability diagram of the posterior probability of Health against the observed share of Health bills (CAP codes) on the validation and non-anchored working bills, for Dawid-Skene, CRE-LCM and the anchored CRE-LCM; point size is the number of bills in the bin

Table 9 CRE-LCM estimates on the bills (priors of Section 3.6; 4 chains of 1500 iterations): prevalence, random-effect standard deviations and variance shares, posterior mean and 95% credible interval, with bulk ESS and $\hat{R}$, without any human code and with 30 anchored bills. The prevalence is reported from the unanchored fit and the variance decomposition from the anchored fit, which converges; the unanchored point estimates of the shares agree with it

Quantity	CRE-LCM, no human code				CRE-LCM, 30 anchored			
	estimate		ESS	$\hat{R}$	estimate		ESS	$\hat{R}$
π_1	0.315	[0.28, 0.35]	1160	1.01	0.311	[0.28, 0.34]	3421	1.01
σ_θ	5.635	[4.95, 6.45]	8	1.42	5.579	[4.99, 6.20]	198	1.02
σ_ϕ	3.066	[2.69, 3.47]	208	1.02	3.120	[2.75, 3.54]	338	1.01
σ_ψ	2.887	[2.58, 3.22]	65	1.05	2.867	[2.58, 3.18]	812	1.01
item share	0.600	[0.52, 0.67]	8	1.46	0.593	[0.53, 0.65]	134	1.03
model share	0.179	[0.14, 0.23]	16	1.18	0.186	[0.15, 0.23]	172	1.02
prompt share	0.159	[0.12, 0.20]	10	1.34	0.157	[0.13, 0.19]	344	1.01
run share	0.063	[0.05, 0.07]	12	1.25	0.063	[0.05, 0.07]	349	1.02




Fig. 6 Posterior predictive check: share of bills by the number of configurations (of 45) labelling the bill Health, observed (bars) and replicated from the fitted CRE-LCM (circles, solid line) and from Dawid-Skene (triangles, dashed line); square-root vertical scale

Table 10 Posterior labels against the CAP human codes: accuracy, Brier score, sensitivity and specificity of the labels (threshold 0.5) and the estimated prevalence, on the 200 validation bills (never anchored) and on the 570 non-anchored working bills

Method	Set	Acc.	Brier	Sens.	Spec.	$\hat{\pi}_1$	π_1 (CAP)
Majority vote	validation	0.970	0.030	0.970	0.970	0.340	0.330
Majority vote	working, non-anchored	0.954	0.046	0.952	0.955	0.309	0.291
Dawid-Skene	validation	0.960	0.040	1.000	0.940	0.370	0.330
Dawid-Skene	working, non-anchored	0.956	0.044	0.976	0.948	0.321	0.291
CRE-LCM	validation	0.970	0.031	1.000	0.955	0.361	0.330
CRE-LCM	working, non-anchored	0.960	0.039	0.946	0.965	0.300	0.291
CRE-LCM anchored	validation	0.970	0.030	1.000	0.955	0.359	0.330
CRE-LCM anchored	working, non-anchored	0.960	0.038	0.940	0.968	0.297	0.291

Relation to bounds and to design-based correction. [Chen et al. \(2026\)](#) show that if the dependence among LLM annotators is left unrestricted, only bounds on prevalence are available. Our results locate the price of that generality: with the dependence structured by the annotation design, point estimates are recovered, at the cost of an assumption (crossed Gaussian effects) that we check by misspecification. Design-based methods ([Egami et al. 2024](#); [Angelopoulos et al. 2023](#)) reach validity from the opposite direction, by spending a labelled subsample on a bias correction (doubly robust in the design-based supervised learning of [Egami et al. \(2024\)](#), a “rectifier” in prediction-powered inference); our anchoring variant shows that a subsample of the same order also removes the estimation difficulties of the latent class model,

Table 11 Sensitivity and specificity of each model–prompt configuration (averaged over its three runs): against the CAP human codes, as estimated by Dawid–Skene, and as implied by CRE-LCM (marginal over the item effects, posterior mean and 95% interval)

Model	Prompt	Sensitivity				Specificity			
		CAP	DS	CRE-LCM		CAP	DS	CRE-LCM	
anthropic	1	0.92	0.92	0.88	[.85, .91]	0.95	0.99	0.98	[.96, .99]
anthropic	2	0.91	0.90	0.94	[.91, .95]	0.96	1.00	0.99	[.98, .99]
anthropic	3	0.77	0.74	0.75	[.71, .79]	0.98	1.00	0.98	[.97, .99]
anthropic	4	0.72	0.73	0.81	[.76, .85]	0.95	0.99	0.98	[.96, .99]
anthropic	5	0.92	0.93	0.85	[.82, .88]	0.93	0.98	0.97	[.96, .98]
gemini	1	0.98	0.98	0.98	[.97, .99]	0.95	1.00	0.99	[.98, .99]
gemini	2	0.96	0.95	0.99	[.98, 1.00]	0.96	1.00	0.99	[.99, 1.00]
gemini	3	0.97	0.98	0.95	[.92, .97]	0.95	1.00	0.99	[.98, 1.00]
gemini	4	0.98	0.99	0.96	[.94, .98]	0.95	1.00	0.99	[.98, .99]
gemini	5	0.98	1.00	0.98	[.96, .99]	0.94	0.99	0.98	[.97, .99]
ollama	1	0.81	0.78	0.82	[.78, .86]	0.97	0.99	0.98	[.97, .99]
ollama	2	0.95	0.92	0.90	[.86, .93]	0.97	1.00	0.99	[.99, 1.00]
ollama	3	0.65	0.65	0.67	[.62, .71]	0.97	0.99	0.99	[.98, .99]
ollama	4	0.80	0.78	0.73	[.69, .78]	0.97	0.99	0.98	[.97, .99]
ollama	5	0.73	0.72	0.79	[.74, .83]	0.97	1.00	0.98	[.97, .99]

and the two approaches are complementary rather than competing—one supplies a measurement model, the other a downstream-inference correction.

Accuracy is not enough. Carlson and Burbano (2026) found that a configuration with 87.6% first-stage accuracy (Claude 3 Opus with chain-of-thought prompting) produced a negative point estimate for a coefficient whose benchmark sign is positive, while one with 55.6% accuracy recovered the positive association, and attributed such reversals to first-stage error whose direction is correlated with the outcome. Our model does not solve that problem, but it makes its ingredients visible: the class-specific model effects a_{mk} show which configurations err asymmetrically, and the posterior class probabilities carry the residual uncertainty into any downstream analysis instead of discarding it.

Limitations. The main model is binary; the multiclass extension is sketched and left for future work. The random effects are Gaussian and crossed; our misspecification arm shows that t_4 item effects leave prevalence, accuracy and the posterior labels unaffected and bias only the item share (Section 4.5), but other departures (bimodal ambiguity, model-by-prompt interactions) are untested. The shared item scale is shrunk by 10–20% at $N = 300$ when θ_i is sampled, which propagates into the item share; anchoring reduces the shrinkage, and with θ_i integrated out the item share is instead slightly overestimated, by 0.04 at $N = 150$ and 0.01 at $N = 600$ (Section 4.4). In the application the configurations are nearly unanimous, and the likelihood is then nearly flat in the scale of the latent effects (quasi-complete separation): the absolute variance shares depend on the prior, while their ranking, the prevalence and the posterior labels do not (Appendix B); when configurations are near-unanimous we recommend reporting the ranking of the shares, with their intervals, rather than their absolute values. A bias shared by every configuration—in the application, the LLMs’ broader notion of health than the CAP coding rule—cannot be separated from the latent class without a gold standard; a small anchored subset that includes boundary items is the remedy (Section 5.3). Two further limitations come from the simulation study. Mean sensitivity is underestimated in small samples, by up to 0.09 at $N = 150$ and 0.01 to 0.05

at $N \geq 300$: the minority class contributes few items, and the ordered constraint and prior shrink μ_1 towards zero; the bias is conservative, smaller than the Dawid–Skene overstatement in the opposite direction, and vanishes with N , so we recommend at least 300 items after adjusting for the size of the minority class. The merged-class mode recurs in 8% of replications at the strongest dependence level, where the likelihood is sharpest; the two-chain diagnostic of Section 3.6 detects it and anchoring removes it. Estimation by MCMC takes minutes per fit; an EM algorithm with quadrature would make the method routine. Finally, LLM versions change: the annotations we report are tied to the version strings and dates in the supplementary material, and the model’s parameters should be understood as properties of those versions, not of the model families.

7 Conclusion

An LLM annotation ensemble is a designed experiment with three crossed factors, and its errors inherit that design. Treating each configuration as an independent rater, as Dawid–Skene does, overstates how good the configurations are and, as dependence grows, biases what they measure. Giving the latent class model the structure of the design yields point estimates without human labels—identification is supported empirically by the simulation study rather than proved—decomposes annotation noise into components that map onto concrete design decisions, and produces class probabilities that are better calibrated in simulation than majority vote or Dawid–Skene. In the application, where the ensemble is near-unanimous, all aggregators label about equally well; the gain lies in the prevalence, whose bias is halved, and in the per-configuration accuracy estimates. The main practical obstacles—a merged-class posterior mode and slow mixing on the prevalence—are addressed by initialising at the Dawid–Skene solution, by integrating the shared item effect out by quadrature, or by anchoring a small labelled subset. Code, prompts, model versions and all annotations are provided for reuse.

Supplementary information. Stan code for all model variants, the full per-cell simulation results and the annotation data are provided as Online Resource 1. The code repository is available for review through an anonymised mirror: <https://anonymous.4open.science/r/cre-lcm-adac-F89F/>.

Appendix A Stan code

A.1 Main model: shared item effect integrated by Gauss–Hermite quadrature

The program used for the simulation study (Section 3.6); the nodes and log-weights of the standard-normal Gauss–Hermite rule are passed as data.

```
// crossed_lcre_q_quad.stan -- variant (q)
// As crossed_lcre.stan, but the shared item-ambiguity effect theta_i is
// INTEGRATED OUT by Gauss-Hermite quadrature inside the class mixture, so the
// N strongly coupled latent parameters theta_i are no longer sampled:
// L_i = sum_k pi_k sum_q w_q prod_{m,p} Binomial(S_imp | R, expit(mu_k + a_mk + b_pk
// + s_th * x_q + phi_im + psi_ip))
```

```

// x_q, w_q are nodes and weights of a standard-normal Gauss-Hermite rule
// (passed as data). The crossed effects phi_im and psi_ip remain sampled
// (non-centred). Removes the slow theta/class random walk seen in the grid.

data {
  int<lower=1> N;
  int<lower=1> M;
  int<lower=1> P;
  int<lower=1> R;
  array[N, M, P] int<lower=0, upper=R> S;
  int<lower=1> Q; // quadrature nodes
  vector[Q] x_q; // nodes (standard normal scale)
  vector[Q] log_w_q; // log weights (sum of weights = 1)
}

transformed data {
  int MP = M * P;
  array[N, MP] int Sv; // S flattened per item, index (m-1)*P + p
  for (i in 1:N) for (m in 1:M) for (p in 1:P) Sv[i, (m - 1) * P + p] = S[i, m, p];
}

parameters {
  real<lower=0, upper=1> pi1;
  ordered[2] mu;
  matrix[M - 1, 2] a_raw;
  matrix[P, 2] b_z;
  vector<lower=0>[2] tau_b;
  real<lower=0> s_th;
  real<lower=0> s_ph;
  real<lower=0> s_ps;
  matrix[N, M] ph_z;
  matrix[N, P] ps_z;
}

transformed parameters {
  matrix[M, 2] a;
  matrix[P, 2] b;
  a[1] = rep_row_vector(0, 2);
  if (M > 1) a[2:M] = a_raw;
  for (k in 1:2) b[, k] = tau_b[k] * b_z[, k];
}

model {
  pi1 ~ beta(1, 1);
  mu ~ normal(0, 2);
  to_vector(a_raw) ~ normal(0, 1);
  to_vector(b_z) ~ std_normal();
  tau_b ~ normal(0, 0.5);
  s_th ~ normal(0, 1);
  s_ph ~ normal(0, 1);
  s_ps ~ normal(0, 1);
  to_vector(ph_z) ~ std_normal();
  to_vector(ps_z) ~ std_normal();

  {
    // class-specific fixed part, flattened over (m, p)
    matrix[MP, 2] fixed;
    for (m in 1:M) for (p in 1:P) for (k in 1:2)
      fixed[(m - 1) * P + p, k] = mu[k] + a[m, k] + b[p, k];
    vector[2] log_pi = [log1m(pi1), log(pi1)]';
    for (i in 1:N) {
      vector[MP] u; // crossed effects for item i
      for (m in 1:M) for (p in 1:P)
        u[(m - 1) * P + p] = s_ph * ph_z[i, m] + s_ps * ps_z[i, p];
      vector[2 * Q] lp;
      for (k in 1:2) for (q in 1:Q)
        lp[(k - 1) * Q + q] = log_pi[k] + log_w_q[q]
          + binomial_logit_lpmf(Sv[i] | R, fixed[, k] + u + s_th * x_q[q]);
    }
  }
}

```

```

        target += log_sum_exp(lp);
    }
}
}

generated quantities {
    vector[N] post1;
    vector[N] log_lik;
    real v_res = square(pi()) / 3;
    real v_tot = square(s_th) + square(s_ph) + square(s_ps) + v_res;
    real share_item = square(s_th) / v_tot;
    real share_model = square(s_ph) / v_tot;
    real share_prompt = square(s_ps) / v_tot;
    real share_run = v_res / v_tot;
    {
        matrix[MP, 2] fixed;
        for (m in 1:M) for (p in 1:P) for (k in 1:2)
            fixed[(m - 1) * P + p, k] = mu[k] + a[m, k] + b[p, k];
        vector[2] log_pi = [log1m(pi1), log(pi1)]';
        for (i in 1:N) {
            vector[MP] u;
            for (m in 1:M) for (p in 1:P)
                u[(m - 1) * P + p] = s_ph * ph_z[i, m] + s_ps * ps_z[i, p];
            vector[2] lk;
            for (k in 1:2) {
                vector[Q] lq;
                for (q in 1:Q)
                    lq[q] = log_w_q[q] + binomial_logit_lpmf(Sv[i] | R, fixed[, k] + u + s_th * x_q[q]);
                lk[k] = log_pi[k] + log_sum_exp(lq);
            }
            log_lik[i] = log_sum_exp(lk);
            post1[i] = exp(lk[2] - log_lik[i]);
        }
    }
}

```

A.2 Variant with sampled item effect

The program with θ_i sampled, used for the identifiability comparison of Section 3.4 (variants (a)–(e) modify its prior, parameterisation, data or initialisation as described there) and for the $\delta \leq 0.5$ cells of the grid.

```

// crossed_lcre.stan
// Binary latent class model with crossed random effects for
// LLM annotation ensembles (model x prompt x run).
// The discrete class c_i is marginalised out.

data {
    int<lower=1> N; // items
    int<lower=1> M; // LLM models
    int<lower=1> P; // prompts
    int<lower=1> R; // runs per (model, prompt)
    array[N, M, P] int<lower=0, upper=R> S; // # runs labelled 1
}

parameters {
    real<lower=0, upper=1> pi1; // prevalence of class 1
    ordered[2] mu; // class intercepts (label identification)
    matrix[M - 1, 2] a_raw; // model effects, model 1 = reference
    matrix[P, 2] b_z; // prompt effects (non-centred)
    vector<lower=0>[2] tau_b;
    real<lower=0> s_th; // item ambiguity SD
    real<lower=0> s_ph; // item x model SD
    real<lower=0> s_ps; // item x prompt SD
    vector[N] th_z;
}

```

```

matrix[N, M] ph_z;
matrix[N, P] ps_z;
}

transformed parameters {
  matrix[M, 2] a;
  matrix[P, 2] b;
  a[1] = rep_row_vector(0, 2);
  if (M > 1) a[2:M] = a_raw;
  for (k in 1:2) b[, k] = tau_b[k] * b_z[, k];
}

model {
  // priors
  pi1 ~ beta(1, 1);
  mu ~ normal(0, 2);
  to_vector(a_raw) ~ normal(0, 1);
  to_vector(b_z) ~ std_normal();
  tau_b ~ normal(0, 0.5);
  s_th ~ normal(0, 1);
  s_ph ~ normal(0, 1);
  s_ps ~ normal(0, 1);
  th_z ~ std_normal();
  to_vector(ph_z) ~ std_normal();
  to_vector(ps_z) ~ std_normal();

  // likelihood, marginalising the latent class
  for (i in 1:N) {
    vector[2] lp = [log1m(pi1), log(pi1)]';
    for (m in 1:M) {
      for (p in 1:P) {
        real u = s_th * th_z[i] + s_ph * ph_z[i, m] + s_ps * ps_z[i, p];
        for (k in 1:2)
          lp[k] += binomial_logit_lpmf(S[i, m, p] | R, mu[k] + a[m, k] + b[p, k] + u);
      }
    }
    target += log_sum_exp(lp);
  }
}

generated quantities {
  vector[N] post1; // P(c_i = 1 | data)
  vector[N] log_lik; // pointwise (per item) log-likelihood for L00
  // variance shares on the latent logistic scale (run-level residual = pi^2/3)
  real v_res = square(pi()) / 3;
  real v_tot = square(s_th) + square(s_ph) + square(s_ps) + v_res;
  real share_item = square(s_th) / v_tot;
  real share_model = square(s_ph) / v_tot;
  real share_prompt = square(s_ps) / v_tot;
  real share_run = v_res / v_tot;
  for (i in 1:N) {
    vector[2] lp = [log1m(pi1), log(pi1)]';
    for (m in 1:M) for (p in 1:P) {
      real u = s_th * th_z[i] + s_ph * ph_z[i, m] + s_ps * ps_z[i, p];
      for (k in 1:2)
        lp[k] += binomial_logit_lpmf(S[i, m, p] | R, mu[k] + a[m, k] + b[p, k] + u);
    }
    log_lik[i] = log_sum_exp(lp);
    post1[i] = exp(lp[2] - log_lik[i]);
  }
}

```

Appendix B Prior sensitivity of the application fits

Table B1 compares the application fits under the priors of Section 3.6 (the primary specification) with a refit under wider priors. When almost every configuration gives the same label, the likelihood is nearly flat in the scale of the latent effects (quasi-complete separation): larger intercepts and random-effect standard deviations fit the near-unanimous counts almost equally well, so their absolute values, and with them the absolute variance shares, are set by the prior. Under the wide priors the intercepts and σ_θ grow, the chains mix poorly and the item share rises at the expense of the run share. The quantities used in Section 5.3 do not depend on this choice: the ranking of the shares, the prevalence, the accuracy and Brier score of the posterior labels against the CAP codes, the per-configuration sensitivities and specificities (which differ by at most 0.025 and 0.005) and the posterior predictive distribution of the agreement counts.

Table B1 Prior sensitivity of the application fits. Original priors (the simulation-study priors of Section 3.6: $N(0, 2)$ on μ , half-normal(0, 1) on the three standard deviations, half-normal(0, 0.5) on τ_k ; four chains of 1500 iterations) against wide priors ($N(0, 10)$ on μ , half-normal(0, 5) on the standard deviations and on τ_k ; four chains of 4000 iterations), without human codes and with 30 anchored bills: posterior mean [95% interval] and bulk ESS; largest $\hat{R}$ over the global parameters; accuracy and Brier score of the posterior labels against the CAP codes on the 570 non-anchored working bills and the 200 validation bills

Quantity	No human code				30 anchored			
	original	ESS	wide	ESS	original	ESS	wide	ESS
π_1	.31 [.28,.35]	1160	.31 [.28,.35]	54	.31 [.28,.34]	3421	.32 [.28,.35]	625
μ_0	-14.5 [-16.2,-12.6]	603	-21.6 [-26.5,-17.8]	13	-14.2 [-15.9,-12.4]	861	-21.7 [-25.4,-18.2]	42
μ_1	7.1 [5.5,8.7]	82	10.9 [6.3,15.3]	19	7.2 [5.5,8.8]	99	11.0 [6.8,15.1]	29
σ_θ	5.6 [4.9,6.5]	8	8.5 [6.9,11.0]	6	5.6 [5.0,6.2]	198	7.8 [6.6,9.1]	10
σ_ϕ	3.1 [2.7,3.5]	208	4.0 [3.4,4.6]	111	3.1 [2.7,3.5]	338	4.1 [3.5,4.7]	44
σ_ψ	2.9 [2.6,3.2]	65	3.5 [3.1,4.0]	139	2.9 [2.6,3.2]	812	3.7 [3.2,4.1]	54
item share	.60 [.52,.67]	8	.69 [.61,.79]	6	.59 [.53,.65]	134	.64 [.56,.72]	7
model share	.18 [.14,.23]	16	.15 [.11,.21]	7	.19 [.15,.23]	172	.18 [.13,.23]	9
prompt share	.16 [.12,.20]	10	.12 [.08,.16]	6	.16 [.13,.19]	344	.14 [.11,.18]	9
run share	.06 [.05,.07]	12	.03 [.02,.04]	7	.06 [.05,.07]	349	.04 [.03,.04]	19
max $\hat{R}$	1.46		1.91		1.03		1.60	
acc., working	0.960		0.961		0.960		0.958	
Brier, working	0.039		0.038		0.038		0.038	
acc., validation	0.970		0.970		0.970		0.970	
Brier, validation	0.031		0.032		0.030		0.032	

References

Atil, B., Aykent, S., Chittams, A., Fu, L., Passonneau, R.J., Radcliffe, E., Rajagopal, G.R., Sloan, A., Tudrej, T., Ture, F., Wu, Z., Xu, L., Baldwin, B.: Non-determinism of “deterministic” LLM system settings in hosted environments. In: Proceedings of the 5th Workshop on Evaluation and Comparison of NLP Systems (Eval4NLP 2025), pp. 135–148. Association for Computational Linguistics, Mumbai, India (2025). <https://doi.org/10.18653/v1/2025.eval4nlp-1.12>

- Angelopoulos, A.N., Bates, S., Fannjiang, C., Jordan, M.I., Zrnic, T.: Prediction-powered inference. *Science* **382**(6671), 669–674 (2023) <https://doi.org/10.1126/science.adi6000>
- Bavaresco, A., Bernardi, R., Bertolazzi, L., Elliott, D., Fernández, R., Gatt, A., Ghaleb, E., Giulianelli, M., Hanna, M., Koller, A., Martins, A.F.T., Mondorf, P., Neplenbroek, V., Pezzelle, S., Plank, B., Schlangen, D., Suglia, A., Surikuchi, A.K., Takmaz, E., Testoni, A.: LLMs instead of human judges? A large scale empirical study across 20 NLP evaluation tasks. In: *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pp. 238–255. Association for Computational Linguistics, Vienna, Austria (2025). <https://doi.org/10.18653/v1/2025.acl-short.20>
- Beath, K.J.: randomLCA: An R package for latent class with random effects analysis. *J. Stat. Softw.* **81**(13), 1–25 (2017) <https://doi.org/10.18637/jss.v081.i13>
- Barrie, C., Palmer, A., Spirling, A.: Replication for language models: Problems, principles, and best practice for political science. Working paper, 2025 (2025). https://arthurspirling.org/documents/BarriePalmerSpirling_TrustMeBro.pdf
- Barrie, C., Palaiologou, E., Törnberg, P.: Prompt stability scoring for text annotation with large language models. arXiv:2407.02039. Version 3, 15 May 2026 (2024). <https://doi.org/10.48550/arXiv.2407.02039>
- Baumann, J., Röttger, P., Urman, A., Wendsjö, A., Plaza-del-Arco, F.M., Gruber, J.B., Hovy, D.: Large language model hacking: Quantifying the hidden risks of using LLMs for text annotation. arXiv:2509.08825. Version 2, 6 October 2025 (2025). <https://doi.org/10.48550/arXiv.2509.08825>
- Carlson, N.A., Burbano, V.: The use of LLMs to annotate data in management research: Foundational guidelines and warnings. *Strateg. Manag. J.* **47**(3), 699–725 (2026) <https://doi.org/10.1002/smj.70023>
- Carpenter, B., Gelman, A., Hoffman, M.D., Lee, D., Goodrich, B., Betancourt, M., Brubaker, M., Guo, J., Li, P., Riddell, A.: Stan: A probabilistic programming language. *J. Stat. Softw.* **76**(1), 1–32 (2017) <https://doi.org/10.18637/jss.v076.i01>
- Camuffo, A., Gambardella, A., Kazemi, S., Malachowski, J., Pandey, A.: Variance-aware LLM annotation for strategy research: Sources, diagnostics, and a protocol for reliable measurement. arXiv:2601.02370. Version 3, 19 January 2026 (2026). <https://doi.org/10.48550/arXiv.2601.02370>
- Chen, X., Rambachan, A., Tamer, E.: Partial identification from LLM prompts. arXiv:2606.15031. Version 2, 24 June 2026 (2026). <https://doi.org/10.48550/arXiv.2606.15031>
- Dawid, A.P., Skene, A.M.: Maximum likelihood estimation of observer error-rates

- using the EM algorithm. *J. R. Stat. Soc. Ser. C Appl. Stat.* **28**(1), 20–28 (1979) <https://doi.org/10.2307/2346806>
- Egami, N., Hinck, M., Stewart, B.M., Wei, H.: Using large language model annotations for the social sciences: A general framework of using predicted variables in downstream analyses. Working paper, version 17 November 2024 (2024). https://naokiegami.com/paper/dsl_ss.pdf
- Gilardi, F., Alizadeh, M., Kubli, M.: ChatGPT outperforms crowd workers for text-annotation tasks. *Proc. Natl. Acad. Sci. U.S.A.* **120**(30), 2305016120 (2023) <https://doi.org/10.1073/pnas.2305016120>
- Jones, B.D., Baumgartner, F.R., Theriault, S.M., Epp, D.A., Khandekar, S., Little, D.: Policy Agendas Project: Codebook. Comparative Agendas Project (2025). <https://www.comparativeagendas.net/pages/master-codebook>
- Jones, G., Johnson, W.O., Hanson, T.E., Christensen, R.: Identifiability of models for multiple diagnostic testing in the absence of a gold standard. *Biometrics* **66**(3), 855–863 (2010) <https://doi.org/10.1111/j.1541-0420.2009.01330.x>
- Liu, J.: What is actually being annotated? Inter-prompt reliability as a measurement problem in LLM-based social science labeling. arXiv:2604.16413 (2026). <https://doi.org/10.48550/arXiv.2604.16413>
- Marbac, M., Biernacki, C., Vandewalle, V.: Latent class model with conditional dependency per modes to cluster categorical data. *Adv. Data Anal. Classif.* **10**(2), 183–207 (2016) <https://doi.org/10.1007/s11634-016-0250-1>
- Messing, S.: Hidden measurement error in LLM pipelines distorts annotation, evaluation, and benchmarking. arXiv:2604.11581. Version 6, 13 May 2026 (2026). <https://doi.org/10.48550/arXiv.2604.11581>
- Nguyen, C., Do, T.-T., Carneiro, G.: Towards the identifiability in noisy label learning: A multinomial mixture modelling approach. arXiv:2301.01405. Version 3, 24 September 2025 (2023). <https://doi.org/10.48550/arXiv.2301.01405>
- Norman, J.D., Rivera, M.U., Hughes, D.A.: Reliability without validity: A systematic, large-scale evaluation of LLM-as-a-judge models across agreement, consistency, and bias. arXiv:2606.19544. Version 1, 17 June 2026 (2026). <https://doi.org/10.48550/arXiv.2606.19544>
- Oberski, D.L.: Beyond the number of classes: Separating substantive from non-substantive dependence in latent class analysis. *Adv. Data Anal. Classif.* **10**(2), 171–182 (2016) <https://doi.org/10.1007/s11634-015-0211-0>
- Oberski, D.L., Kollenburg, G.H., Vermunt, J.K.: A Monte Carlo evaluation of three methods to detect local dependence in binary data latent class models. *Adv. Data*

- Anal. Classif. **7**(3), 267–279 (2013) <https://doi.org/10.1007/s11634-013-0146-2>
- Qu, Y., Tan, M., Kutner, M.H.: Random effects models in latent class analysis for evaluating accuracy of diagnostic tests. *Biometrics* **52**(3), 797–810 (1996) <https://doi.org/10.2307/2533043>
- Raykar, V.C., Yu, S.: Eliminating spammers and ranking annotators for crowdsourced labeling tasks. *J. Mach. Learn. Res.* **13**, 491–518 (2012)
- Salter-Townshend, M., Murphy, T.B.: Mixtures of biased sentiment analysers. *Adv. Data Anal. Classif.* **8**(1), 85–103 (2014) <https://doi.org/10.1007/s11634-013-0150-6>
- Vacek, P.M.: The effect of conditional dependence on the evaluation of diagnostic tests. *Biometrics* **41**(4), 959–968 (1985) <https://doi.org/10.2307/2530967>
- Vehtari, A., Gelman, A., Gabry, J.: Practical Bayesian model evaluation using leave-one-out cross-validation and WAIC. *Stat. Comput.* **27**(5), 1413–1432 (2017) <https://doi.org/10.1007/s11222-016-9696-4>
- Vehtari, A., Gelman, A., Simpson, D., Carpenter, B., Bürkner, P.-C.: Rank-normalization, folding, and localization: An improved $\hat{R}$ for assessing convergence of MCMC (with discussion). *Bayesian Anal.* **16**(2), 667–718 (2021) <https://doi.org/10.1214/20-BA1221>
- Wilkerson, J., Adler, E.S., Jones, B.D., Baumgartner, F.R., Freedman, G., Theriault, S.M., Craig, A., Epp, D.A., Sullivan, M.E., Khandekar, S., Little, D.: Policy Agendas Project: Congressional Bills. Comparative Agendas Project, dataset version 19.3 (2025). <https://www.comparativeagendas.net/us>
- Warrens, M.J.: Inequalities between multi-rater kappas. *Adv. Data Anal. Classif.* **4**(4), 271–286 (2010) <https://doi.org/10.1007/s11634-010-0073-4>
- Warrens, M.J.: Cohen’s weighted kappa with additive weights. *Adv. Data Anal. Classif.* **7**(1), 41–55 (2013) <https://doi.org/10.1007/s11634-013-0123-9>
- Whitehill, J., Ruvolo, P., Wu, T., Bergsma, J., Movellan, J.: Whose vote should count more: Optimal integration of labels from labelers of unknown expertise. In: *Advances in Neural Information Processing Systems* 22, pp. 2035–2043 (2009)
- Xie, Y., Chen, Z., Albert, P.S.: A crossed random effects modeling approach for estimating diagnostic accuracy from ordinal ratings without a gold standard. *Stat. Med.* **32**(20), 3472–3485 (2013) <https://doi.org/10.1002/sim.5784>
- Xu, Z., Khatri, V., Dai, Y., Liu, X., Li, S., Zhang, X., Yu, R.: Enhancing LLM-based data annotation with error decomposition. In: *Proceedings of the 16th International Learning Analytics and Knowledge Conference (LAK26)*, pp. 325–335. ACM, Bergen, Norway (2026). <https://doi.org/10.1145/3785022.3785070>

- Xu, H., Li, X., Shen, C., Hui, S.L., Grannis, S.: Incorporating conditional dependence in latent class models for probabilistic record linkage: Does it matter? *Ann. Appl. Stat.* **13**(3), 1753–1790 (2019) <https://doi.org/10.1214/19-AOAS1256>
- Zhang, Y., Martinez, I.: From stochasticity to signal: A Bayesian latent state model for reliable measurement with LLMs. arXiv:2510.23874. Version 4, 23 April 2026 (2025). <https://doi.org/10.48550/arXiv.2510.23874>